

Computer Networks

Jiaqi Zheng

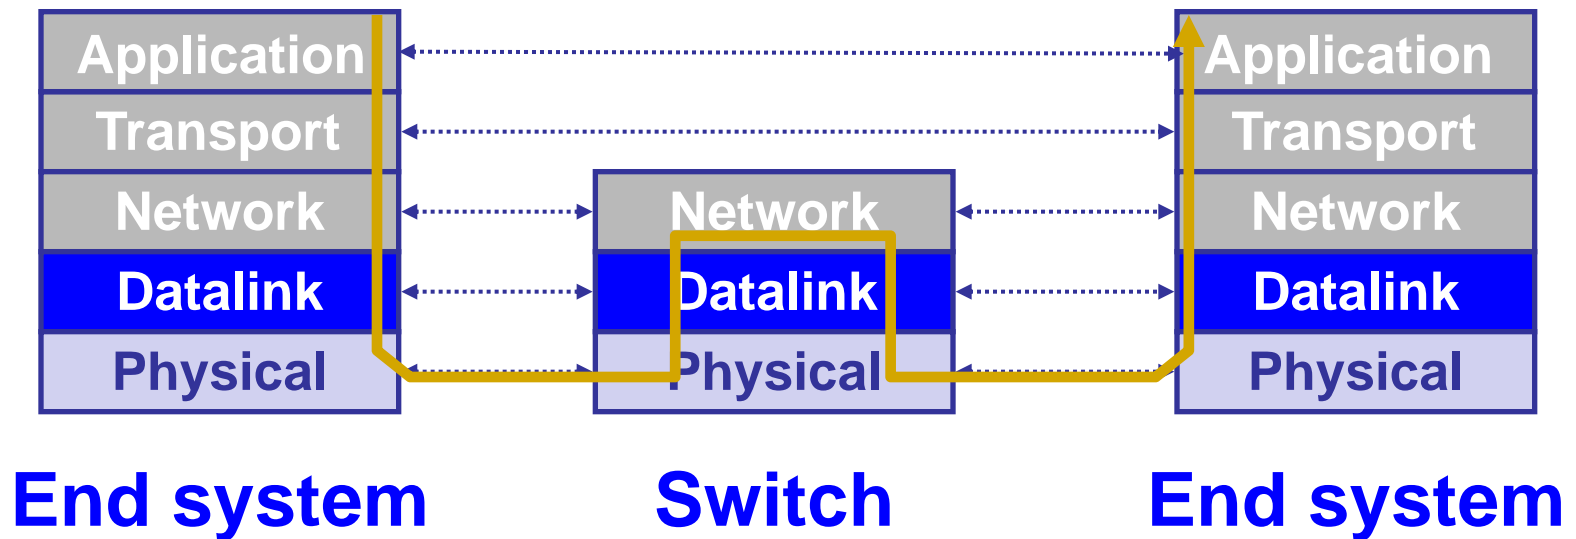
Material with thanks Mosharaf Chowdhury, and many other colleagues.

Agenda

- Data link layer

Data link layer


- Present everywhere
- Transfers data between adjacent nodes or between nodes on the same local area network



Data link layer

- Provides four primary services
 - Framing
 - » Encapsulates network layer data
 - Link access
 - » Medium access control (MAC) protocol defines when to transmit frames
 - Reliable delivery
 - » Primarily for mediums with high error rates (e.g., wireless)
 - Error detection and correction

Packets are now “frames”



- Frames encapsulate network layer packets
- Link layer protocols are implemented in h/w
- Frame formats can change based on link layer protocol

Point-to-point vs. broadcast medium

- **Point-to-point**: dedicated pairwise communication
 - E.g., long-distance fiber link
 - E.g., Point-to-point link b/n Ethernet switch and host
- **Broadcast**: shared wire or medium
 - Traditional Ethernet (pre ~2000)
 - 802.11 wireless LAN

Multiple access algorithm

- Context: a shared broadcast channel
 - Must avoid having multiple nodes speaking at once
 - » Otherwise, collisions lead to garbled data
 - Need distributed algorithm to determine which node can transmit
- Three classes of techniques
 - **Channel partitioning**: divide channel into pieces
 - **Taking turns**: scheme for deciding who transmits
 - **Random access**: allow collisions, and then recover
 - » More in the Internet style!

Random access MAC protocols

- When node has packet to send
 - Transmit at full channel data rate **w/o** coordination
- Two or more transmitting nodes \Rightarrow collision
 - Data lost
- Random access MAC protocol specifies
 - How to detect and recover from collisions
- Examples
 - ALOHA and Slotted ALOHA
 - CSMA, CSMA/CD, CSMA/CA (wireless)

Ethernet


- Invented as a broadcast technology
 - Hosts share channel
 - Each packet received by all attached hosts
 - CSMA/CD for media access control
- Modern Ethernets are “switched” (later)
 - Point-to-point links between switches and between a host and switch
 - No sharing \Rightarrow no CSMA/CD
 - » Uses “self learning” and “spanning tree” algorithms for routing

CSMA (Carrier Sense Multiple Access)

- CSMA: listen before transmit
 - If channel sensed idle: transmit entire frame
 - If channel sensed busy, defer transmission
- Human analogy: don't interrupt others!
- Does not eliminate all collisions
 - Why?

CSMA collisions

- **Propagation delay**: two nodes may not hear each other before sending
- CSMA reduces but does not eliminate collisions
- Collision: entire packet transmission time wasted
 - Distance and propagation delay affect collision probability




CSMA/CD (Collision Detection)

- CSMA/CD: carrier sensing, deferral as in CSMA
 - Collisions detected within short time
 - Colliding transmissions aborted, reducing wastage
- Collision detection easy in wired (broadcast) LANs
 - Compare transmitted, received signals
- Collision detection difficult in wireless LANs
 - Later!

CSMA/CD (Collision Detection)

- For this to work, need restrictions on minimum frame size and maximum distance
 - Why?



Limits on CSMA/CD network length



- Latency depends on physical length of link
 - Time to propagate a frame from one end to other
- Suppose A sends a frame at time t
 - And B sees an idle line at a time just before $t + d$
 - ... so B happily starts transmitting a frame
- B detects a collision, and sends jamming signal
 - But A cannot see collision until $t + 2d$

Limits on CSMA/CD network length



- A needs to wait for time **$2d$** to detect collision
 - So, A should keep transmitting during this period
 - AND keep an eye out for a possible collision
- Imposes restrictions; e.g., for 10 Mbps Ethernet
 - **Maximum length** of the wire: 2,500 meters
 - **Minimum length** of a frame: 512 bits (64 bytes)

Three key ideas of random access

- Carrier sense

- Listen before speaking and don't interrupt
- Checking if someone else is already sending data
- ... and waiting till the other node is done

- Collision detection

- If someone else starts talking at the same time, stop
 - Make sure everyone knows there was a collision!
- Realizing when two nodes are transmitting at once
- ...by detecting that the data on the wire is garbled

Three key ideas of random access

- Randomness
 - Don't start talking again right away
 - Waiting for a random time before trying again

How long should you wait?

- Should it be immediate?
- Should it be a random number with a fixed distribution?

Ethernet: CSMA/CD Protocol

- Carrier sense: wait for link to be idle
- Collision detection: listen while transmitting
 - No collision: transmission is complete
 - Collision: abort transmission & send jam signal
- Random access: **binary exponential back-off**
 - After collision, wait a random time before retrying
 - After m^{th} collision, choose K randomly from $\{0, \dots, 2^m - 1\}$
 - » Wait for $K * 512$ bit times before trying again
 - » If transmission occurring when ready to send, wait until end of transmission (CSMA)

Efficiency of CSMA/CD

- Efficiency is defined as the long-run fraction of time during which frames are being transmitted without collision
- d_{prop} = max propagation time between two adapters
- d_{trans} = time to transmit a max-sized frame

$$\text{Efficiency} \approx \frac{1}{1 + 5 d_{\text{prop}} / d_{\text{trans}}}$$

Efficiency of CSMA/CD

- $d_{\text{prop}} \rightarrow 0$
 - Efficiency approaches 1
 - Colliding nodes abort immediately
- $d_{\text{trans}} \rightarrow \infty$
 - Efficiency approaches 1
 - Each frames uses the channel for a long time

$$\text{Efficiency} \approx \frac{d_{\text{trans}}}{d_{\text{trans}} + 5 d_{\text{prop}}}$$


SWITCHED ETHERNET

Broadcast vs. switched Ethernet

- Invented as a broadcast technology
 - Hosts share channel
 - Each packet received by all attached hosts
 - CSMA/CD for media access control
- Modern Ethernets are “switched”
 - Point-to-point links between switches and between a host and switch
 - No sharing \Rightarrow no CSMA/CD
 - » Uses “self learning” and “spanning tree” algorithms for routing

Why switched Ethernet?

- Enables concurrent communication
 - Host A can talk to C, while B talks to D
 - No collisions and no need for CSMA/CD
 - No constraints on link lengths, etc.



The evolution of Ethernet


- Changed almost everything except the frame format
 - From the shared media coax cables to dedicated links
 - From 3 Mbit/s to 100 Gbit/s
 - From electrical signaling to optical
- **Lesson:** the right interface can accommodate many changes
 - Evolve the implementation while maintaining the interface (backward compatibility)

Topics

- Frames and framing
- Addressing
- Routing
- Forwarding
- Discovery

Ethernet “Frames”

- Encapsulates IP datagram



- **Preamble**: 7 bytes for clock synchronization and 1 byte to indicate start of frame
- **Addresses**: 6 bytes
- **Type**: 2 bytes, higher-layer protocol (e.g., IP)
- **Data payload**: max 1500 bytes, min 46 bytes
- **CRC**: 4 bytes for error detection

Framing frames


- Physical layer puts bits on a link
- But, two hosts connected on the same physical medium need to be able to exchange frames
 - Service provided by the link layer
 - Implemented by the network adaptor
- **Framing problem**: how does the link layer determine where each frame begins and ends?

Simple approach: Count bytes

- Sender includes number of bytes in header



- Receiver extracts this number of bytes of body
- What if the Count field is corrupted?




- L2 will frame the wrong bytes → a framing error
- CRC tells you to discard this frame, but what about the next one?

Desynchronization

- Once framing on a link is desynchronized, it can stay that way
- Need a method to **resynchronize**


Framing with sentinel bits

- Delineate frame with special “sentinel” bit pattern
 - e.g., 01111110 \Rightarrow start, 01111111 \Rightarrow end



- What if sentinel occurs within frame?
- Solution: bit stuffing
 - Sender always inserts a 0 after five 1s in the frame contents
 - Receiver always removes a 0 appearing after five 1s

When receiver sees five 1s...



- If next bit 0, remove it; begin counting again
 - Because this must be a stuffed bit; we can't be at beginning/end of frame (those had six or seven 1s)
- If next bit 1 (i.e., we've seen six 1s) then:
 - If following bit is 0, this is start of frame
 - » Because the receiver has seen 01111110
 - If following bit is 1, this is end of frame
 - » Because the receiver has seen 01111111

Example: sentinel bits

- Original data, including start/end of frame:
 - 01111110011111101111101111100101111111
- Sender rule: five 1s → insert a 0
 - After bit stuffing at the sender:
 - 01111110011111010111110011111000101111111
- Receiver rule: five 1s and next bit 0 → remove 0
 - 01111110011111101111101111100101111111

Topics

- Frames and framing
- Addressing
- Routing
- Forwarding
- Discovery

Medium Access Control (MAC) Address

- MAC address
 - Numerical address associated with a network adapter
 - Flat name space of 48 bits (e.g., 00-15-C5-49-04-A9 in HEX)
 - Unique, hard-coded in the adapter when it is built
- Hierarchical Allocation
 - **Blocks**: assigned to vendors (e.g., Dell) by the IEEE
 - » First 24 bits (e.g., 00-15-C5-**-**-**)
 - **Adapter**: assigned by the vendor from its block
 - » Last 24 bits

MAC address vs. IP address

MAC Addresses

- Hard-coded when adapter is built
- Flat name space of 48 bits (e.g., 00-0E-9B-6E-49-76)
- Like a social security number
- Portable, and can stay the same as the host moves
- Used to get packet between interfaces on same network


IP Addresses

- Configured, or learned dynamically
- Hierarchical name space of 32 bits (e.g., 12.178.66.9)
- Like a postal mailing address
- Not portable, and depends on where the host is attached
- Used to get a packet to destination IP subnet


Topics

- Frames and framing
- Addressing
- Routing
- Forwarding
- Discovery

Routing with switched Ethernet?




“Routing” with broadcast Ethernet



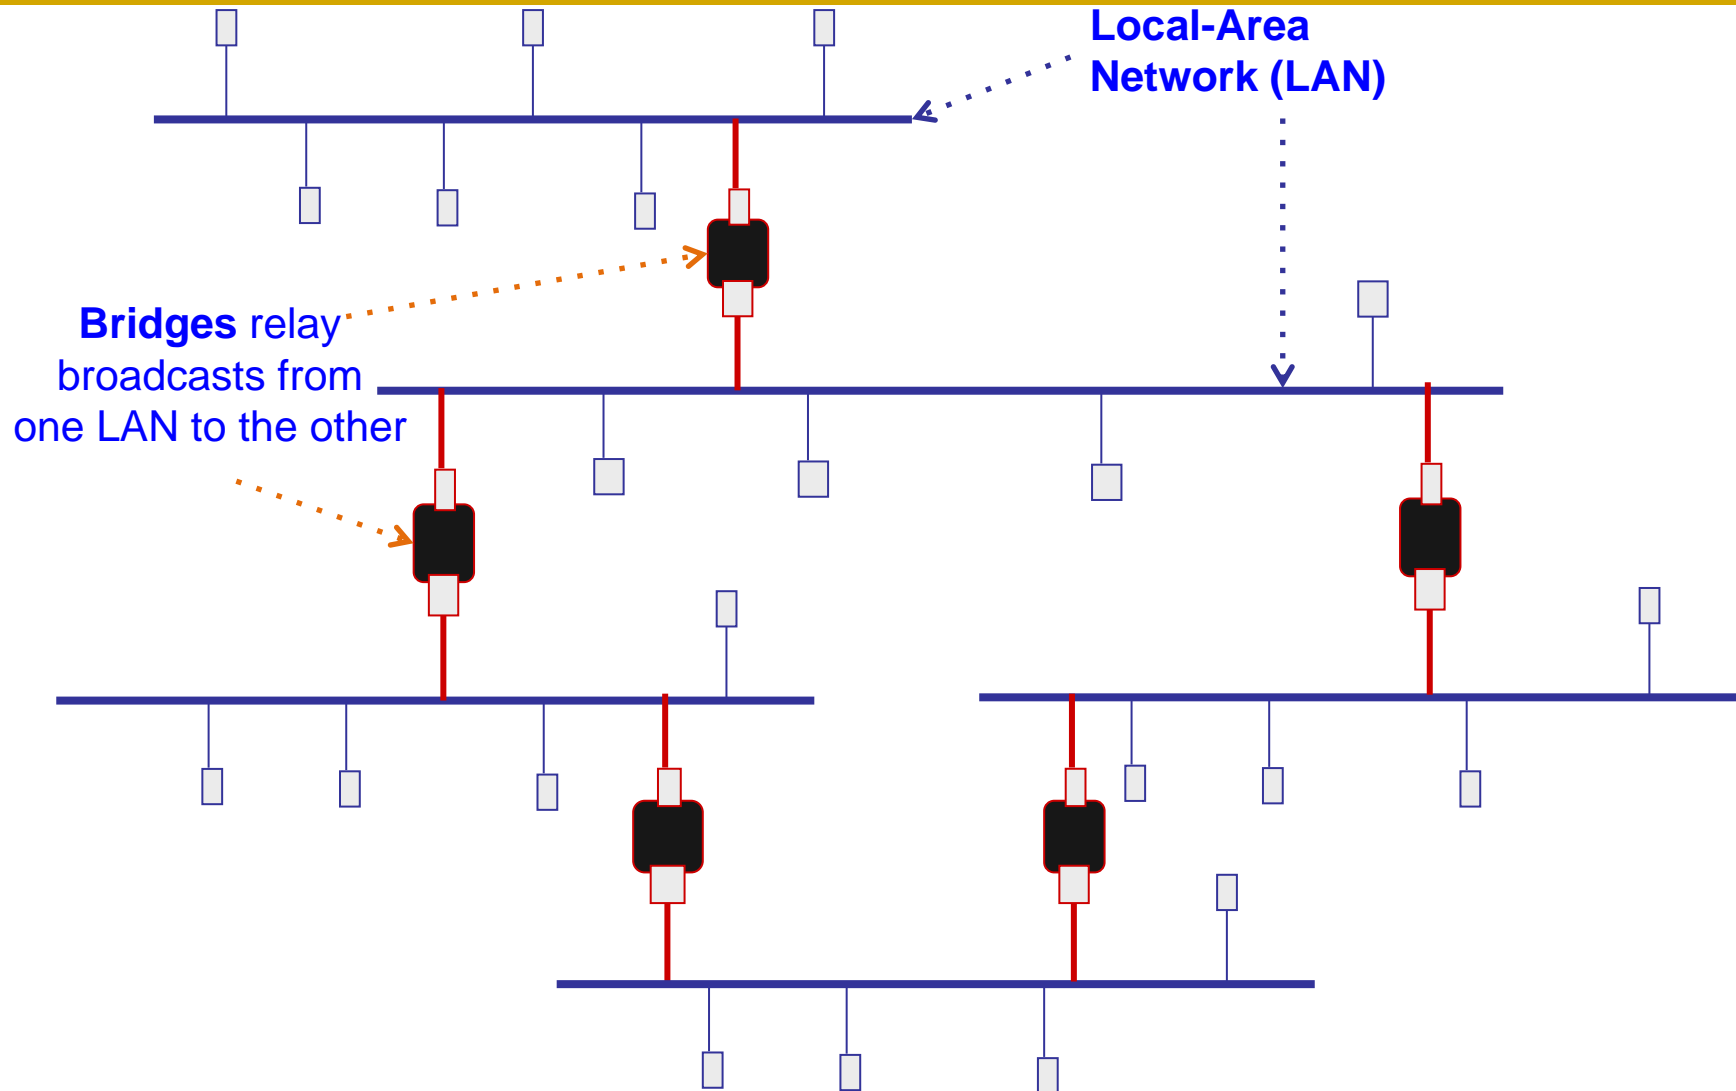
- Sender transmits frame onto broadcast link
- Each receiver's link layer passes the frame to the network layer:
 - If destination address matches the receiver's MAC address OR if the destination address is the broadcast MAC address (ff:ff:ff:ff:ff:ff)

“Routing” with broadcast Ethernet

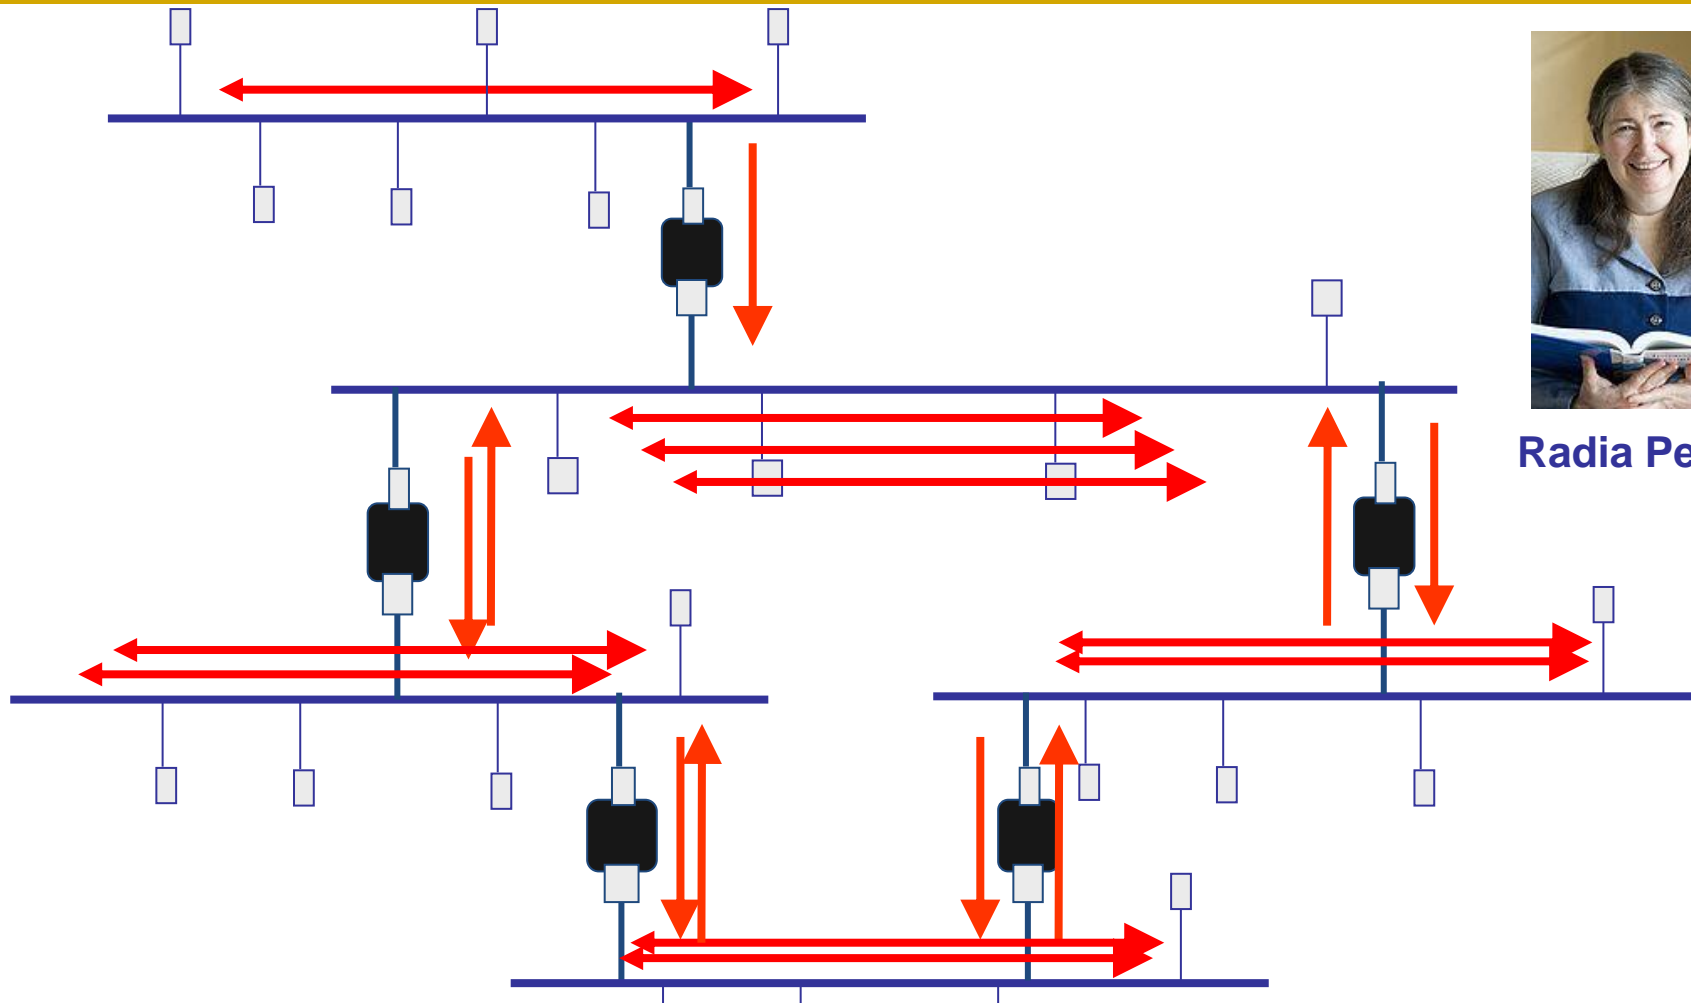


- Ethernet is “plug-n-play”
- A new host plugs into the Ethernet and is good to go
 - No configuration by users or network operators
 - Broadcast as a means of bootstrapping comm.

Routing in extended LANs



The “broadcast storm” problem




Radia Perlman

Perlman's idea: eliminate loops in the topology


Easiest way to avoid loops

- Use a topology where loops are impossible!
- Take arbitrary topology and build a **spanning tree**
 - Sub-graph that includes all vertices but contains no cycles
 - Links not in the spanning tree are not used to forward frames


Consider a graph




A spanning tree



Another spanning tree





Yet another spanning tree



Spanning tree protocol (Perlman'85)

- Protocol by which bridges construct a spanning tree
- Nice properties
 - Zero configuration (by operators or users)
 - Self healing
- Still used today

From extended LANs to switched Ethernet



Switched Ethernet

- Constraints (for backward compatibility)
 - No changes to end-hosts
 - Maintain plug-n-play aspect
- Earlier Ethernet achieved plug-n-play by leveraging a broadcast medium
 - Can we do the same in a switched topology?

Summary

- Data link layer transfers data between adjacent nodes or nodes connected to the same switch
- Ethernet evolved from a broadcast medium to switched



Why does Ethernet not use LS/DV?

- Concerns over scalability
 - Flat MAC addresses cannot be aggregated like IP addresses
- Legacy

Why does Ethernet not use LS/DV?

- Concerns over scalability
 - Flat MAC addresses cannot be aggregated like IP addresses
- Legacy
 - Backward compatibility with broadcast Ethernet
 - Desire to maintain Ethernet's plug-n-play behavior
 - How broadcast Ethernet evolved