

Computer Networks

Jiaqi Zheng

Material with thanks Mosharaf Chowdhury, and many other colleagues.

Agenda

- CDN: Content Distribution Network
- DNS: Domain Name System

Recap:

Improving HTTP performance

- Optimizing connections using three “P”s
 - Persistent connections
 - Parallel/concurrent connections
 - Pipelined transfers over the same connection
- Caching
 - Forward proxy: close to clients
 - Reverse proxy: close to servers
- Replication

Replication

- Replicate popular Websites across many machines
 - Spreads load across servers
 - Places content closer to clients
 - Helps when content isn't cacheable

Content Distribution Networks (CDN)

- Caching and replication as a service
- Large-scale distributed storage infrastructure (usually) administered by one entity
 - e.g., Akamai has servers in 20,000+ locations
- Combination of caching and replication
 - **Pull**: Direct result of clients' requests (caching)
 - **Push**: Expectation of high access rate (replication)
- Can do some processing to handle dynamic webpage content

Cost-effective content delivery

- General theme: multiple sites hosted on shared physical infrastructure
 - Efficiency of statistical multiplexing
 - Economies of scale (volume pricing, etc.)
 - Amortization of human operator costs
- Examples:
 - CDNs
 - Web hosting companies
 - Cloud infrastructure

CDN example – Akamai

- Akamai creates new domain names for each client
 - e.g., a128.g.akamai.net for cnn.com
- The client content provider modifies content so that embedded URLs reference new domains
 - “Akamaize” content
 - e.g., <http://www.cnn.com/image-of-the-day.gif> becomes <http://a128.g.akamai.net/image-of-the-day.gif>
- Requests now sent to CDN’s infrastructure

How to direct clients to particular replicas?

- In order to
 - Balancing load across server replicas
 - Pairing clients with nearby servers to decrease latency and overall bandwidth usage

DNS: DOMAIN NAME SYSTEM

Internet names & addresses

- Machine addresses: e.g., 141.212.113.143
 - Router-usable labels for machines
 - Conforms to network structure (the “where”)
- Machine names: e.g., cse.umich.edu
 - Human-usable labels for machines
 - Conforms to organizational structure (the “who”)
- The Domain Name System (DNS) is how we map from one to the other
 - A [directory](#) service

Why?

- Convenience
 - Easier to remember www.google.com than 216.58.216.100
- Provides a **level of indirection!**
 - Decoupled names from addresses
 - Many uses beyond just naming a specific host

DNS: History

- Initially all host-address mappings were in a `hosts.txt` file (in `/etc/hosts`):
 - Maintained by the Stanford Research Institute (SRI)
 - Changes were submitted to SRI by email
 - New versions of `hosts.txt` periodically FTP'd from SRI

DNS: History (cont'd)

- As the Internet grew this system broke down
 - SRI couldn't handle the load
 - Names were not unique
 - Hosts had inaccurate copies of `hosts.txt`
- The Domain Name System (DNS) was invented to fix this

Goals

- Uniqueness: no naming conflicts
- Scalable
 - Many names and frequent updates (secondary)
- Distributed, autonomous administration
 - Ability to update my own (machines') names
 - Don't have to track everybody's updates
- Highly available
- Lookups are fast
- Perfect consistency is a **non-goal**

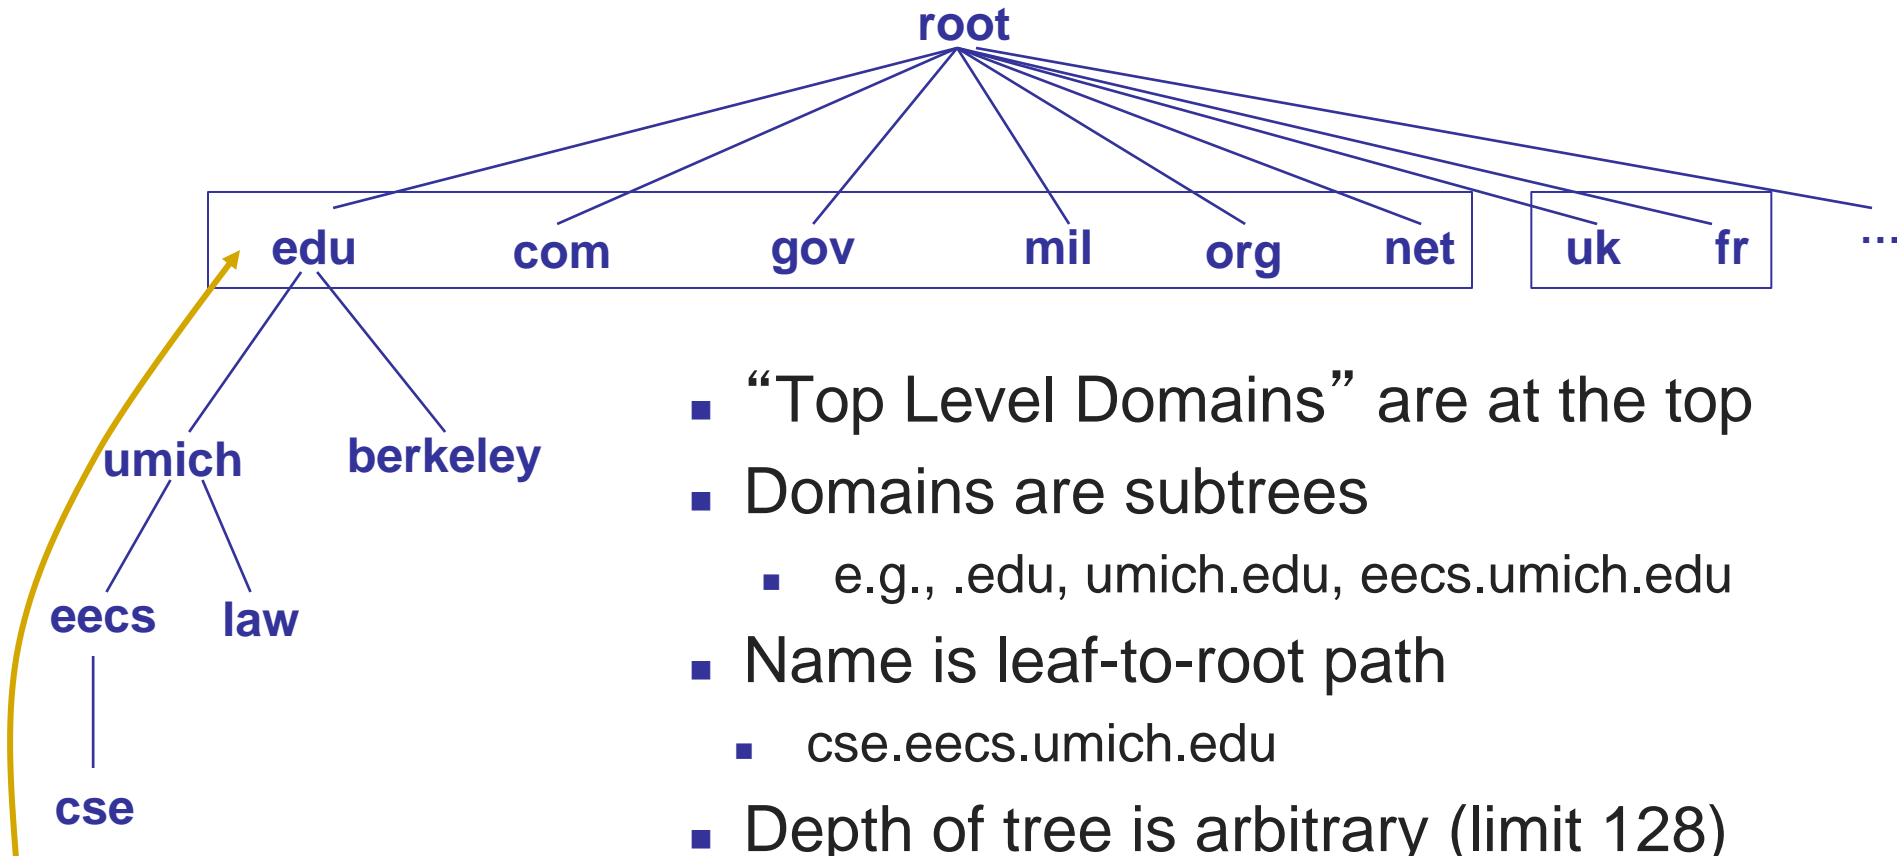
How?

- Partition the namespace
- Distribute administration of each partition
 - Autonomy to update my own (machines') names
 - Don't have to track everybody's updates
- Distribute name resolution for each partition
- How should we partition things?

Key idea: Hierarchy

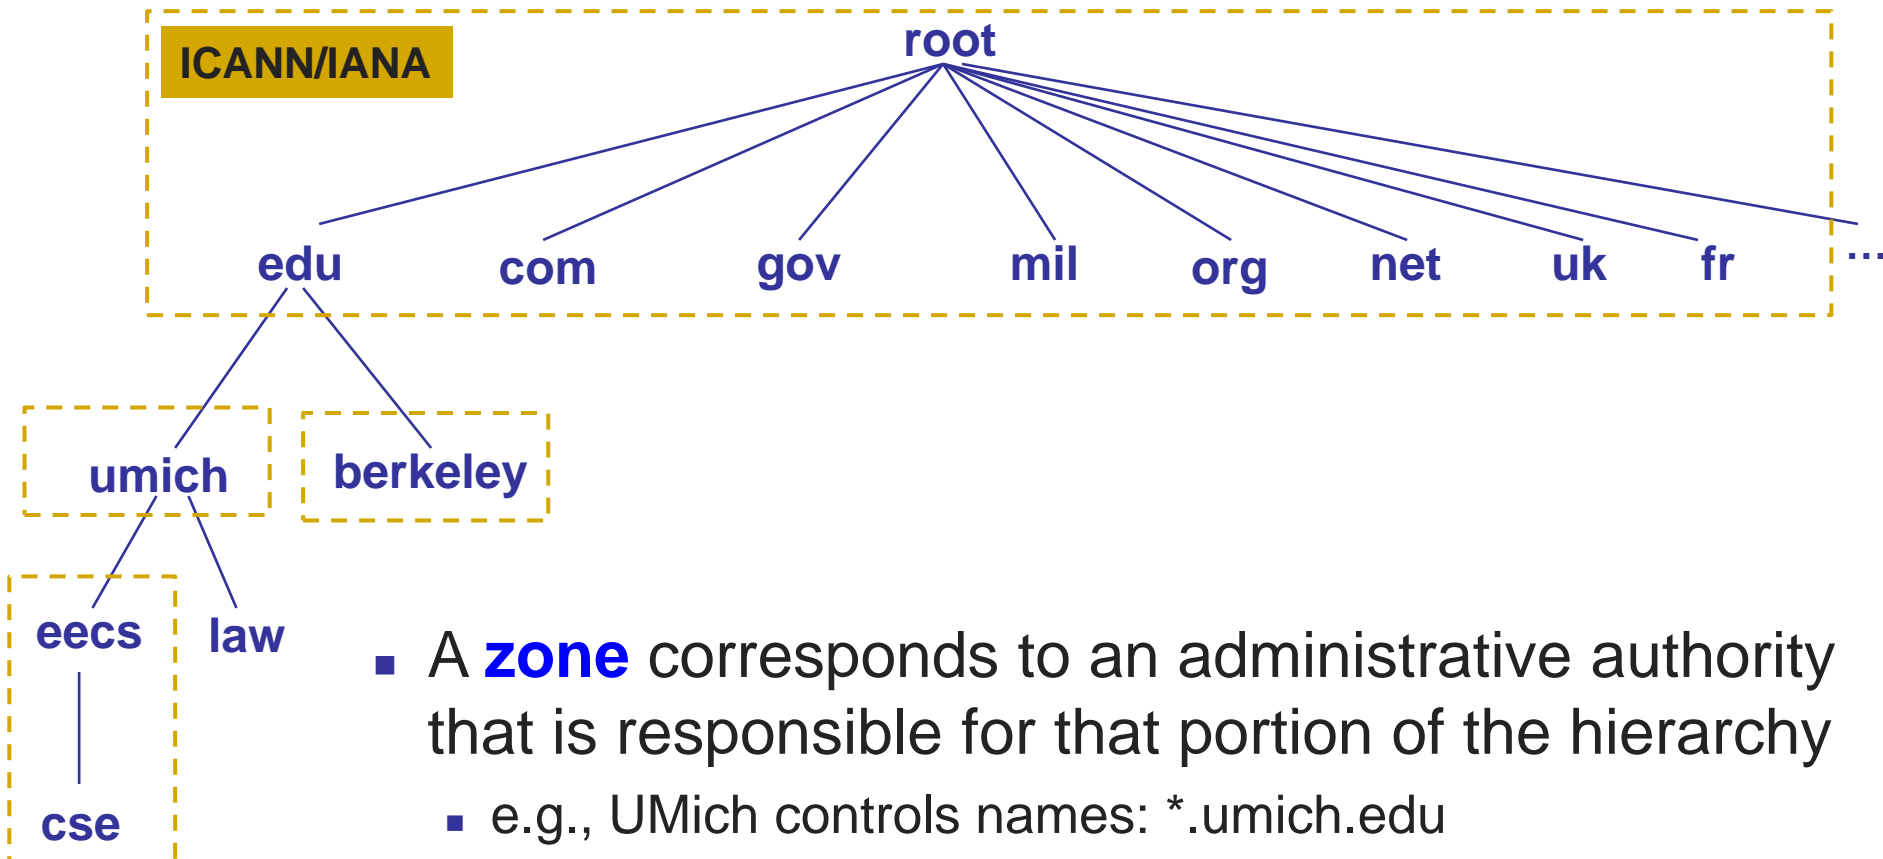
- Three intertwined hierarchies
 - Hierarchical namespace
 - » As opposed to original flat namespace
 - Hierarchically administered
 - » As opposed to centralized
 - (Distributed) hierarchy of servers
 - » As opposed to centralized storage

Hierarchical namespace



- “Top Level Domains” are at the top
- Domains are subtrees
 - e.g., .edu, umich.edu, eeecs.umich.edu
- Name is leaf-to-root path
 - cse.eecs.umich.edu
- Depth of tree is arbitrary (limit 128)
- Name collisions trivially avoided
 - Each domain is responsible

Hierarchical administration



- A **zone** corresponds to an administrative authority that is responsible for that portion of the hierarchy
 - e.g., UMich controls names: *.umich.edu
 - e.g., EECS controls names: *.eecs.umich.edu

Server hierarchy

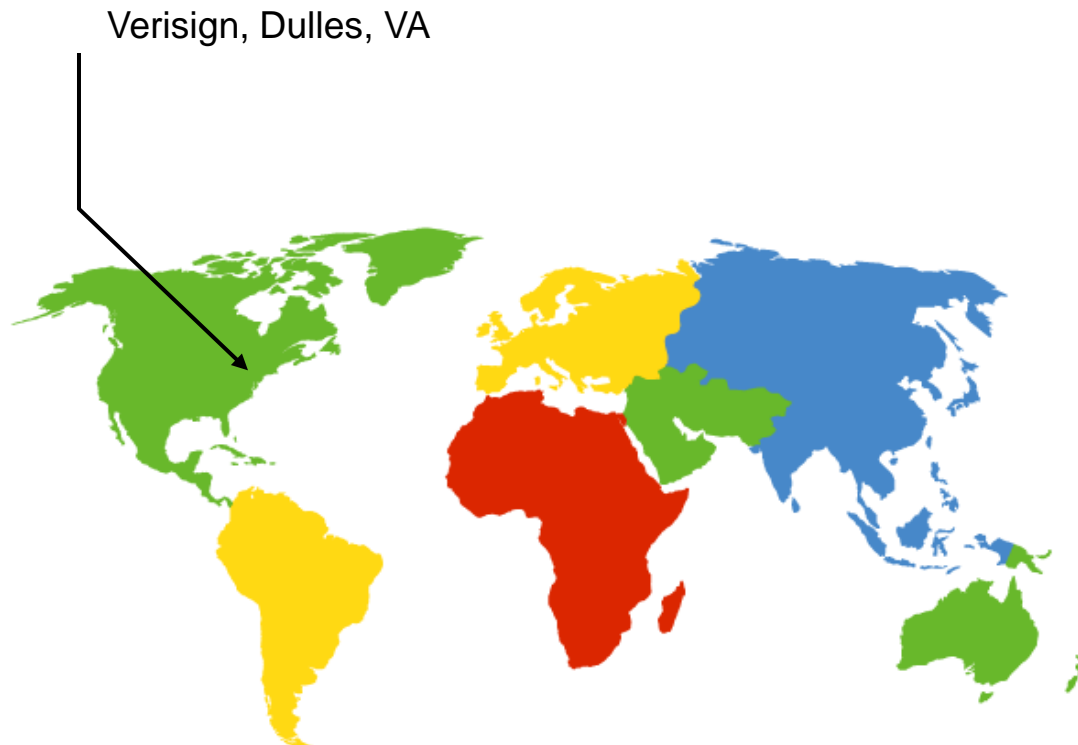
- Top of hierarchy: **Root servers**
 - Location hardwired into other servers
- Next Level: **Top-level domain (TLD) servers**
 - .com, .edu, etc.
 - Managed professionally
- Bottom Level: **Authoritative DNS servers**
 - Actually store the name-to-address mapping
 - Maintained by the corresponding administrative authority

Server hierarchy

- Each server stores a (small!) subset of the total DNS database
- An authoritative DNS server stores “resource records” for all DNS names in the domain that it has authority for
- Each server needs to know other servers that are responsible for the other portions of the hierarchy
 - Every server knows the root
 - Root server knows about all top-level domains

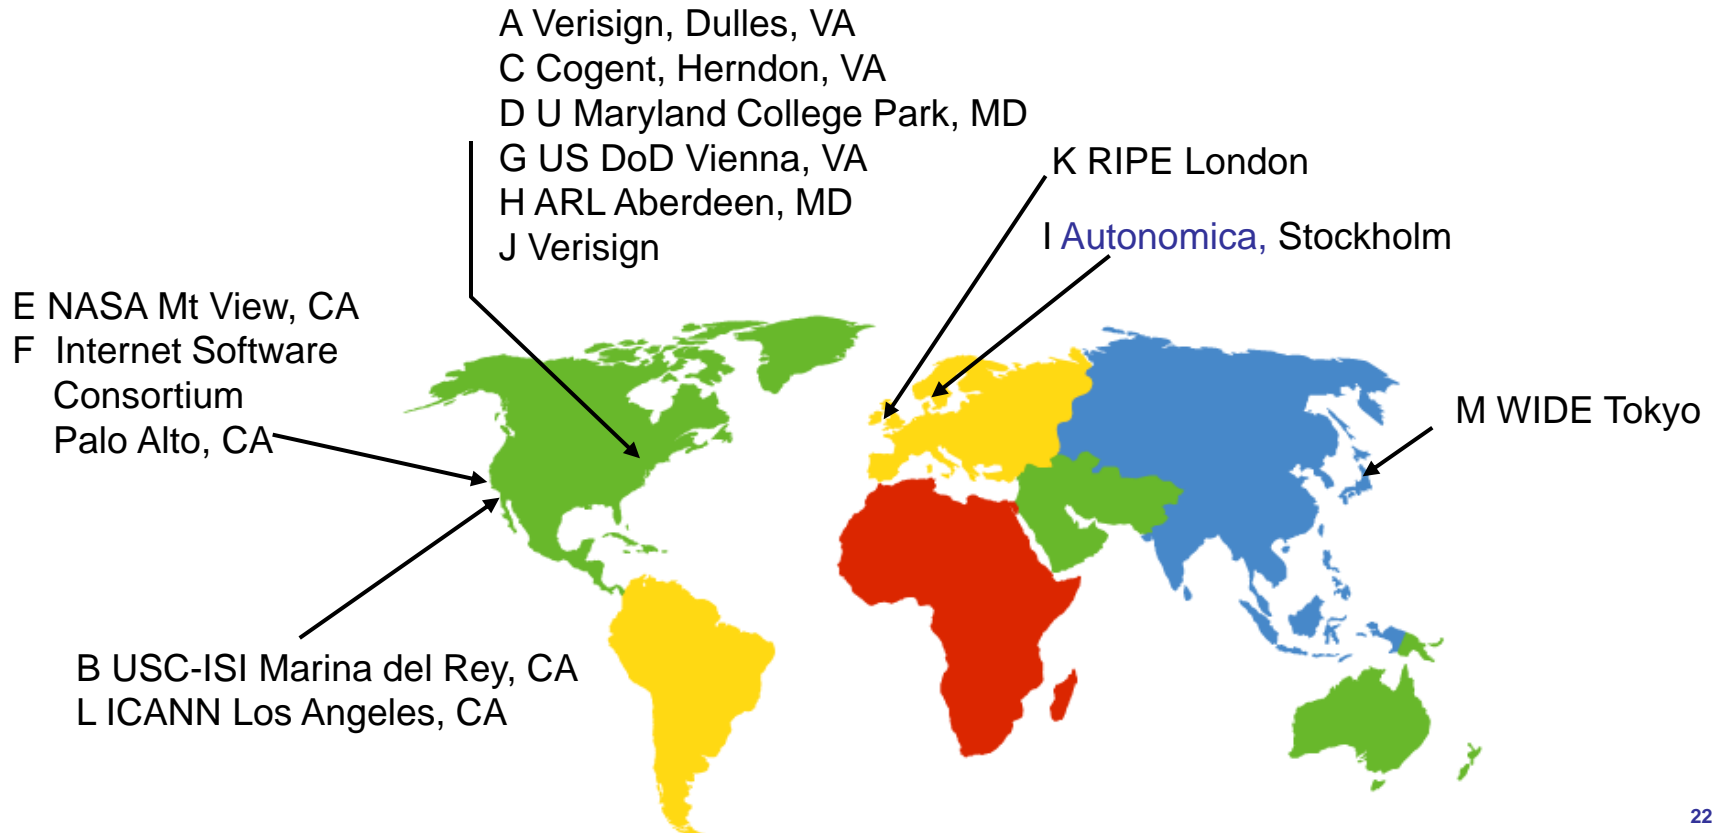
DNS root

- Located in Virginia, USA
- How do we make the root scale?



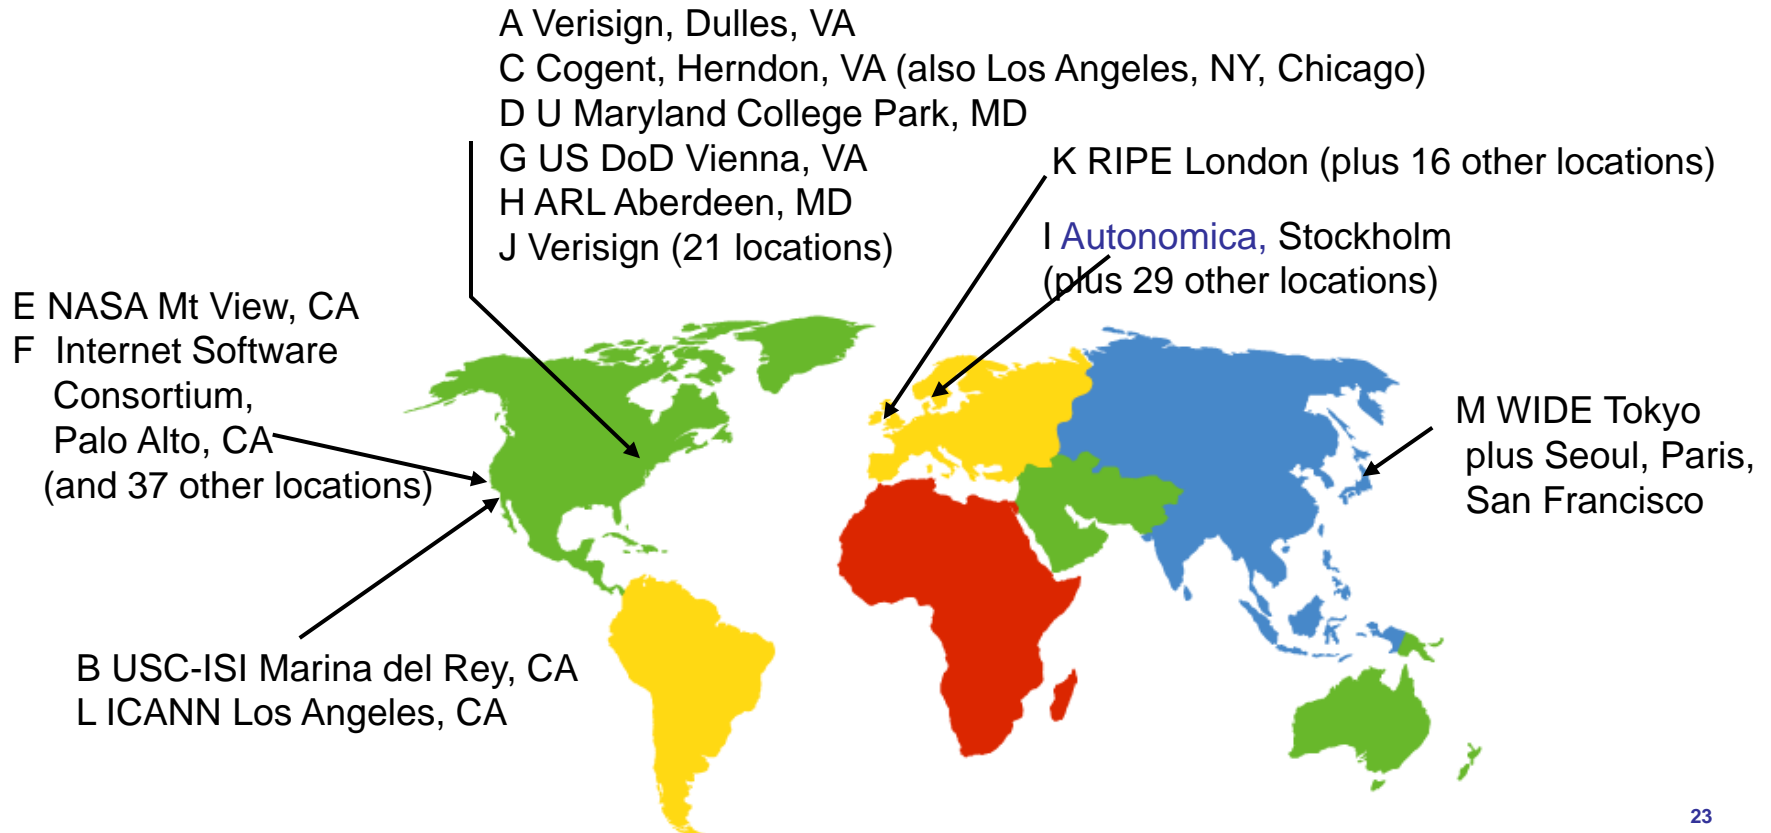
DNS root servers

- 13 root servers (labeled A-M; see <http://www.root-servers.org/>)



DNS root servers

- 13 root servers replicated via **anycast**



Anycast in a nutshell

- Routing finds shortest paths to destination
- If several locations are given the same address, then the network will deliver the packet to the closest location with that address
- Characteristics
 - Very robust
 - Requires no modification to routing algorithms

DNS records

- DNS servers store **resource records (RRs)**
 - RR is (name, value, type, TTL)
- Type = A: (\rightarrow **A**ddress)
 - name = hostname
 - value = IP address
- Type = NS: (\rightarrow **N**ame **S**erver)
 - name = domain
 - value = name of DNS server for domain

DNS records (cont'd)

- Type = CNAME: (→ Canonical Name)
 - name = alias name for some “canonical” (real) name
 - » e.g., cse.umich.edu is really cse.eecs.umich.edu
 - value = canonical name
- Type = MX: (→ Mail eXchanger)
 - name = domain in email address
 - value = name(s) of mail server(s)


Inserting Resource Records into DNS

- Register foobar.com at registrar (GoDaddy)
 - Provide registrar with names and IP addresses of your authoritative name server(s)
 - Registrar inserts RR pairs into the .com TLD server:
 - »(foobar.com, dns1.foobar.com, NS)
 - »(dns1.foobar.com, 212.44.9.129, A)
- Store resource records in your server dns1.foobar.com
 - e.g., type A record for www.foobar.com
 - e.g., type MX record for foobar.com


Using DNS (Client/App View)

- Two components
 - Local DNS servers
 - Resolver software on hosts
- Local DNS server (“default name server”)
 - Clients configured with default server’s address OR learn it via a host configuration protocol (e.g., DHCP)
- Client application
 - Obtain DNS name (e.g., from URL)
 - Do `gethostbyname()` to trigger DNS request to its local DNS server


Name resolution




Name resolution




Name resolution




Name resolution



Name resolution: Recursive



Name resolution: Iterative



Two ways to resolve a name

- Recursive name resolution
 - Ask server to do it for you
- Iterative name resolution
 - Ask server who to ask next
- The iterative example we saw is a mix of both!

DNS protocol

- **Query** and **Reply** messages; both with the same message format
 - Header: identifier, flags, etc.
 - Plus resource records
 - See text/section for details
- Client–server interaction on UDP Port 53
 - Spec supports TCP too, but not always implemented

Goals: Are we there yet?

- Uniqueness: No naming conflicts
- Scalable
- Distributed, autonomous administration
- Highly available?

Reliability

- **Replicated** DNS servers (primary/secondary)
 - Name service available if at least one replica is up
 - Queries can be load-balanced between replicas
- Usually, UDP used for queries
 - Reliability, if needed, must be implemented on UDP
- Try alternate servers on timeout
 - **Exponential backoff** when retrying same server
- Same identifier for all queries
 - Don't care which server responds

Goals: Are we there yet?

- Uniqueness: No naming conflicts
- Scalable
- Distributed, autonomous administration
- Highly available
- Fast lookups?

DNS caching

- Performing all these queries takes time
 - Up to 1-second latency before starting download
- Caching can greatly reduce overhead
 - The top-level servers very rarely change
 - Popular sites (e.g., www.cnn.com) visited often
 - Local DNS server often has the information cached
- How DNS caching works
 - DNS servers cache responses to queries
 - Responses include a “time to live” (TTL) field
 - Server deletes cached entry after TTL expires

Negative caching

- Remember things that don't work
 - Misspellings like `www.cnn.comm` and `www.cnnn.com`
 - These can take a long time to fail the first time
 - Good to remember that they don't work so the failure takes less time the next time around
- Negative caching is optional
 - Not widely implemented

Important properties of DNS

- Administrative delegation and hierarchy enables:
 - Easy unique naming
 - “Fate sharing” for network failures
 - Reasonable trust model
 - Caching increases scalability and performance

DNS provides indirection

- Addresses can change underneath
 - Move `www.cnn.com` to `4.125.91.21`
- Name could map to multiple IP addresses
 - Load-balancing (CDN)
 - Reducing latency by picking nearby servers (CDN)
- Multiple names for the same address
 - E.g., many services (mail, www) on same machine
 - E.g., aliases like `www.cnn.com` and `cnn.com`
- This flexibility applies only within domain!

Summary

- CDNs improve web performance
 - Via replication and caching
 - Good server selection
- DNS allows us to go to webpages without having to memorize IP addresses
 - Allows a level of indirection that enables many functionalities including CDN server selection