

# Computer Networks

Jiaqi Zheng

*Material with thanks Mosharaf Chowdhury, and many other colleagues.*

# Agenda

---

- Software-defined networking

# The field of networking

---

- We have built a great artifact – the Internet
  - It grew mostly unrelated to the academic research, which came later
- CS networking today is largely the study of the Internet
- BUT we do not really have an academic discipline

# Building an artifact, not a discipline

---

- Other fields in “systems”: OS, DB, etc.
  - Teach basic principles
  - Are easily managed
  - Continue to evolve
- Networking:
  - Teach big bag of protocols
  - Notoriously difficult to manage
  - Evolves very slowly
- **Networks are much more primitive and less understood than other computer systems**

# A tale of two planes

---

- **Data plane:** forwarding packets
  - Based on local forwarding state
- **Control plane:** computing that forwarding state
  - Involves coordination with rest of system

# Original goals for the control plane

---

- **Basic connectivity:** route packets to destination
  - Local state computed by routing protocols
  - Globally distributed algorithms
- **Inter-domain policy:** find policy-compliant paths
  - Done by globally distributed BGP
- What other goals are there in running a network?

# Extended roles of the control plane

---

- Performs various network management tasks
  - For example,
    - » Where to route?
    - » How much to route?
    - » At what rate to route?
    - » Should we route at all?
    - » ...


# Traffic engineering

---

- Want to avoid persistent overloads on links
- Choose routes to spread traffic load across links



# Traffic engineering: Difficult




**Q:** What if network operator wants u-to-z traffic to flow along uvwz, x-to-z traffic to flow xwyz?

**A:** Need to define link weights so traffic routing algorithm computes routes accordingly (or need a new routing algorithm)!

Link weights are only control “knobs”


# Traffic engineering: Difficult



**Q:** What if network operator wants to split u-to-z traffic along uvwz and uxyz (load balancing)?

**A:** Can't do it (or need a new routing algorithm)

# Traffic engineering: Difficult



**Q:** What if w wants to route the two flows differently?

**A:** Can't do it (with LS or DV)

# Network management has many goals

---

- Achieving these goals is job of the control plane...
- ...which currently involves many mechanisms
- **Globally distributed:** Routing algorithms
- **Manual/scripted configuration:** ACLs
- **Centralized computation:** Traffic engineering

# Bottom line

---

- Many different control plane mechanisms
- Each designed from scratch for their intended goal
- Encompassing a wide variety of implementations
  - Distributed, manual, centralized,...
- None of them particularly well designed
- Network control plane is a complicated mess!

# “The Power of Abstraction”

---

- “Modularity based on abstraction is the way things get done”
  - Barbara Liskov
- Abstractions → Interfaces → Modularity

# Analogy: Mainframe to PC evolution

---

## Vertical integration, closed

- Specialized application
- Specialized operating system
- Specialized hardware

## Open interfaces

- Arbitrary applications
- Commodity operating systems
- Microprocessor

**We want the same for networking!**

# Many control plane mechanisms

---

- Variety of goals, no modularity
  - Routing: distributed routing algorithms
  - Isolation: ACLs, Firewalls,...
  - Traffic engineering: adjusting weights,...
- Control Plane: mechanism without abstraction
  - Too many mechanisms, not enough functionality



# Task: Compute forwarding state

---

- Consistent with low-level hardware/software
  - Which might depend on particular vendor
- Based on entire network topology
  - Because many control decisions depend on topology
- For all routers/switches in network
  - Every router/switch needs forwarding state

# Our current approach

---

- Design one-off mechanisms that solve all three
- Other fields would define abstractions for each subtask
- ...and so should we!

# Separate concerns with abstractions

---

- Be compatible with low-level hardware/software
  - Need an abstraction for general forwarding model
- Make decisions based on entire network
  - Need an abstraction for network state
- Compute configuration of each physical device
  - Need an abstraction that simplifies configuration

# #1: Forwarding abstraction

---

- Express intent independent of implementation
  - Don't want to deal with proprietary HW and SW
- **OpenFlow** is current proposal for forwarding
  - Standardized interface to switch
  - Configuration in terms of flow entries: <header, action>
- Design details concern exact nature of:
  - Header matching
  - Allowed actions

# Two important facets to OpenFlow

---

- Switches accept external control messages
  - Not closed, proprietary boxes
- Standardized flow entry format
  - So switches are interchangeable

# Separate concerns with abstractions

---

- Be compatible with low-level hardware/software
  - Need an abstraction for general forwarding model
- Make decisions based on entire network
  - Need **an abstraction for network state**
- Compute configuration of each physical device
  - Need an abstraction that simplifies configuration

# #2: Network state abstraction

---

- Abstract away various distributed mechanisms
- Abstraction: **global network view**
  - Annotated network graph provided through an API
- Creates a logically centralized view of the network (Network Operating System)
  - Runs on replicated servers in network (“controllers”)
- Information flows both ways
  - Information from routers/switches to form “view”
  - Configurations to routers/switches to control forwarding

# Network Operating System

---

- Think of it as a centralized link-state algorithm
- Switches send connectivity info to controller
- Controller computes forwarding state
  - Some control program that uses the topology as input
- Controller sends forwarding state to switches
- Controller is replicated for resilience
  - System is only “logically centralized”
- Complicated protocols replaced with simple graph algorithms



# Separate concerns with abstractions

---

- Be compatible with low-level hardware/software
  - Need an abstraction for general forwarding model
- Make decisions based on entire network
  - Need an abstraction for network state
- Compute configuration of each physical device
  - Need an abstraction that simplifies configuration

# #3: Specification abstraction

---

- Control mechanism expresses desired behavior
  - Whether it be isolation, access control, or QoS
- It should not be responsible for implementing that behavior on physical network infrastructure
  - Requires configuring the forwarding tables in each switch
- **Abstract view** of network
  - Models only enough detail to specify goals
  - Will depend on task semantics


# Separate concerns with abstractions

---

- Be compatible with low-level hardware/software
  - Forwarding abstraction
- Make decisions based on entire network
  - Network state abstraction
- Compute configuration of each physical device
  - Specification abstraction


# Traditional fully decentralized control plane

- Individual routing algorithm components in every router interact in the control plane



# Logically centralized control plane

- A distinct (typically remote) controller interacts with local control agents (CAs)



# Each goal is an app via specification abstraction

---

- What if an operator wants X?
- What if a customer wants to do weighted traffic splitting?
- ...
- **There is an app for it!**
  - Write your own routing protocol, load balancing algorithm, access control policies

# Reason about each app via network state abstraction

---

- Now that the network is not distributed anymore and is a simple graph, we can **verify** whether whatever we specified...
  - ...makes sense
  - ...likely to work
  - ...likely to work with the rest
- No more *umm-I don't no-maybe*

# Logically centralized control plane

---

- A distinct (typically remote) controller interacts with local control agents (CAs)
- Each router contains a **flow table**
- Each entry of the flow table defines a **match-action** rule
- Entries of the flow table is computed and distributed by the (logically) centralized controller



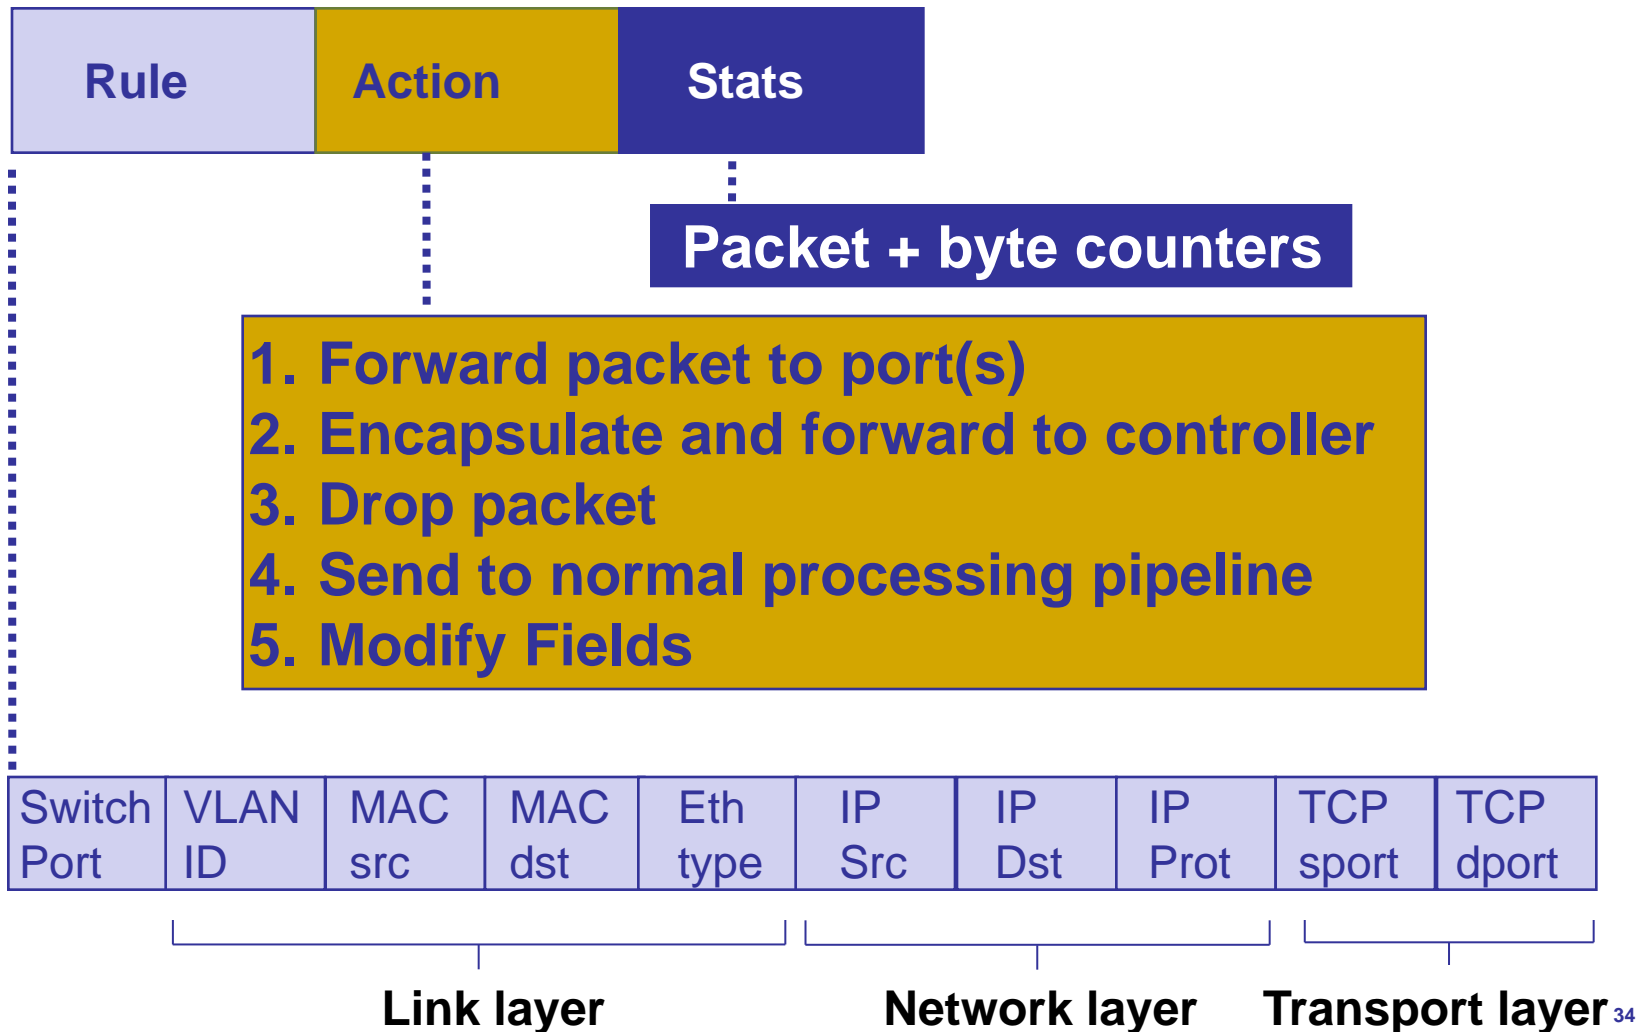
# OpenFlow data plane abstraction

---

- Flow is defined by header fields
- Generalized forwarding: simple packet-handling rules
  - **Pattern**: match values in packet header fields
  - **Actions**: for matched packet: drop, forward, modify, matched packet or send matched packet to controller
  - **Priority**: disambiguate overlapping patterns
  - **Counters**: #bytes and #packets

1. `src=1.2.*.*`, `dest=3.4.5.*` → drop
2. `src = *.*.*`, `dest=3.4.*.*` → forward(2)
3. `src=10.1.2.3`, `dest= *.*.*` → send to controller

# OpenFlow: Flow table entries



# Forwarding abstraction


---

**Match + Action:** unifies different kinds of devices

- Router
  - Match: longest destination IP prefix
  - Action: forward out a link
- Switch
  - Match: destination MAC address
  - Action: forward or flood
- Firewall
  - Match: IP addresses and TCP/UDP port numbers
  - Action: permit or deny
- NAT
  - Match: IP address and port
  - Action: rewrite address and port

# OpenFlow example

match	action
IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(3)




**Example:** datagrams from hosts h5 and h6 should be sent to h3 or h4, via s1 and from there to s2

match	action
ingress port = 1 IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(4)

match	action
ingress port = 2 IP Dst = 10.2.0.3	forward(3)
ingress port = 2 IP Dst = 10.2.0.4	forward(4)

# OpenFlow protocol

## OpenFlow Controller



- Operates between controller, switch
- TCP used to exchange messages
  - Optional encryption
- Three classes of OpenFlow messages:
  - Controller-to-switch
  - Asynchronous (switch to controller)
  - Symmetric (misc.)

# OpenFlow: Controller-to-switch messages

---

- Key controller-to-switch messages
  - **Features**: controller queries switch features, switch replies
  - **Configure**: controller queries/sets switch configuration parameters
  - **Modify-state**: add, delete, modify flow entries in the OpenFlow tables
  - **Packet-out**: controller can send this packet out of specific switch port

# OpenFlow: Switch-to-controller messages

---

- Key switch-to-controller messages
  - **Packet-in**: transfer packet (and its control) to controller. See packet-out message from controller
  - **Flow-removed**: flow table entry deleted at switch
  - **Port status**: inform controller of a change on a port
- Network operators do not “program” switches by creating/sending OpenFlow messages directly.
  - Instead, they use higher-level abstraction at controller

# SDN: Many challenges remain

---

- Hardening the control plane: dependable, reliable, performance-scalable, secure distributed system
  - Robustness to failures: leverage strong theory of reliable distributed system for control plane
  - Dependability, security: “baked in” from day one?
- Networks, protocols meeting mission-specific requirements
  - E.g., real-time, ultra-reliable, ultra-secure
- Internet-scaling



# Some progress for SDN in the wide area

---

- Google and Microsoft use SDN to manage traffic between datacenters
- One centralized controller to rule the entire world (well, their world)

# Google's WAN-SDN (B4)

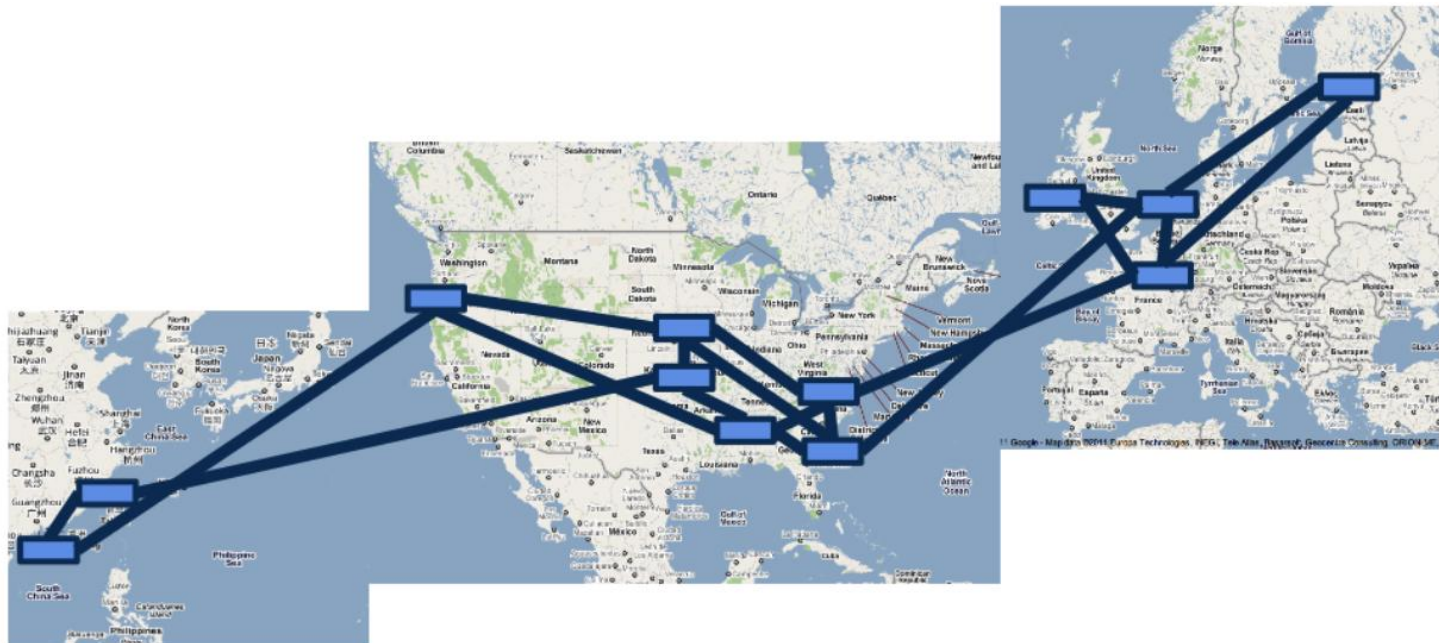


Figure 1: B4 worldwide deployment (2011).

# Software-defined IXP (SDX)

---

- More expressive policies where ASes meet
- <https://noise-lab.net/projects/software-defined-networking/sdx/>

# Summary

---

- Abstractions beget modularity
  - Modularity is (almost always) good